

# Human Activity Recognition

February 3, 2019

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from tabulate import tabulate
```

```
In [151]: plt.style.use('ggplot')
%config InlineBackend.figure_format = 'retina'
sns.set(rc={'figure.figsize':(10,6)})
```

```
In [3]: !gpustat
```

```
instance-gpu  Sat Feb  2 10:51:28 2019
```

```
[0] Tesla P100-PCIE-16GB | 39'C,  0 % |    82 / 16280 MB | root(82M)
```

```
In [4]: !ls
```

```
breast_cancer_1.csv  Human Activity Recognition.ipynb  UCI HAR Dataset
```

```
HAR_RAW_1.csv      nohup.out                        UCI HAR Dataset.names
```

```
In [5]: !ls "UCI HAR Dataset"/"train"/"Inertial Signals"
```

```
body_acc_x_train.txt  body_gyro_x_train.txt  total_acc_x_train.txt
body_acc_y_train.txt  body_gyro_y_train.txt  total_acc_y_train.txt
body_acc_z_train.txt  body_gyro_z_train.txt  total_acc_z_train.txt
```

```
In [6]: %%sh
```

```
head -n 1 "UCI HAR Dataset"/"train"/"Inertial Signals"/body_acc_x_train.txt | wc -w
```

```
128
```

```
In [7]: %%sh
```

```
cat "UCI HAR Dataset"/"train"/y_train.txt | wc -l
```

```
cat "UCI HAR Dataset"/"test"/y_test.txt | wc -l
```

```
7352
```

```
2947
```

### 0.0.1 Preparing the dataset

```
In [8]: SIGNALS = [
        'body_acc_x',
        'body_acc_y',
        'body_acc_z',
        'body_gyro_x',
        'body_gyro_y',
        'body_gyro_z',
        'total_acc_x',
        'total_acc_y',
        'total_acc_z'
    ]

In [9]: def get_data(subset, part, oneHotEncoded=True):
        if part == 'features':
            all_signals = []
            for signal in SIGNALS:
                all_signals.append(np.genfromtxt(
                    "UCI HAR Dataset/{}/Inertial Signals/{_}.txt".format(subset, signal, s
                return np.transpose(all_signals, (1, 2, 0))
        elif part == 'labels':
            y = np.genfromtxt("UCI HAR Dataset/{s}/y_{s}.txt".format(s=subset))
            if oneHotEncoded:
                return pd.get_dummies(y).values
            return y

In [10]: %%time
        X_train = get_data(subset='train', part='features')
        X_test = get_data(subset='test', part='features')

CPU times: user 8.35 s, sys: 626 ms, total: 8.98 s
Wall time: 8.98 s

In [11]: %%time
        y_train = get_data(subset='train', part='labels')
        y_test = get_data(subset='test', part='labels')

CPU times: user 32.2 ms, sys: 3.15 ms, total: 35.3 ms
Wall time: 33.2 ms

In [12]: X_train.shape

Out[12]: (7352, 128, 9)

In [13]: X_test.shape

Out[13]: (2947, 128, 9)
```

```
In [14]: y_train.shape
```

```
Out[14]: (7352, 6)
```

```
In [15]: y_test.shape
```

```
Out[15]: (2947, 6)
```

## 0.0.2 Building a LSTM structure

```
In [16]: X_train.shape[1:]
```

```
Out[16]: (128, 9)
```

```
In [17]: from keras.layers import LSTM, Dropout, Dense
         from keras.models import Sequential
```

Using TensorFlow backend.

```
In [18]: from talos import live
         import talos
```

```
In [23]: def get_best_lstm_model(x_train, y_train, x_val, y_val, params):
```

```
    timesteps, data_point_size = X_train.shape[1:]
    n_classes = y_train.shape[1]

    model = Sequential()

    model.add(LSTM(units=params['first_neuron'],
                    input_shape=(timesteps, data_point_size)))
    model.add(Dropout(rate=params['dropout']))

    model.add(Dense(n_classes, activation='softmax'))

    model.compile(optimizer=params['optimizer'],
                  loss=params['losses'],
                  metrics=['accuracy'])

    history = model.fit(x_train, y_train,
                        validation_data=(x_val, y_val),
                        batch_size=params['batch_size'],
                        callbacks=[live()],
                        epochs=params['epochs'],
                        verbose=2)

    return history, model
```

```
In [24]: y_train.shape[1]
```

```
Out[24]: 6
```

```
In [25]: params = {
    'optimizer': ['Adam'],
    'losses': ['binary_crossentropy'],
    'first_neuron': [16, 32, 64],
    'batch_size': [16],
    'epochs': [20],
    'dropout': [0.3, 0.5, 0.7]
}
```

```
In [148]: get_ipython().run_line_magic('matplotlib', 'inline')
```

### 0.0.3 Randomised GridSearch with one layer LSTM

```
In [ ]: # The output of this cell has been omitted because it
# the connection from the jupyter server is broken due to
# poor internet connections.

# This code runs for almost 6 hours on my GPU (Nvidia Titan XP)
t = talos.Scan(x=X_train, y=y_train,
              x_val=X_test, y_val=y_test,
              model=get_best_lstm_model,
              params=params, search_method='random',
              shuffle=True, dataset_name='HAR_RAW',
              experiment_no='1', seed=42)
```

```
In [31]: t.details
```

```
Out[31]: complete_time      02/02/19/16:10
          experiment_name    HAR_RAW_1
          grid_downsample    None
          random_method      uniform_mersenne
          reduce_loss        False
          reduction_interval  50
          reduction_method    None
          reduction_metric    val_acc
          reduction_threshold 0.2
          reduction_window    20
          x_shape             (7352, 128, 9)
          y_shape             (7352, 6)
          dtype: object
```

```
In [35]: print('We have fit {} number of models'.format(len(t.saved_models)))
```

We have fit 9 number of models

```
In [100]: epoch_results = t.data.drop(['losses', 'optimizer',
                                         'epochs', 'batch_size', 'round_epochs'], axis=1)
```

```

epoch_results['val_loss'] = epoch_results.val_loss.astype(float, inplace=True)
epoch_results['loss'] = epoch_results.loss.astype(float, inplace=True)
epoch_results['val_acc'] = epoch_results.val_acc.astype(float, inplace=True)
epoch_results['acc'] = epoch_results.acc.astype(float, inplace=True)
epoch_results.sort_values(by='val_acc', ascending=False, inplace=True)
columns = ['Val Loss', 'Train Loss',
           'Val Accuracy', 'Train Accuracy',
           'LSTM cells', 'Dropout']

```

```

In [101]: print(tabulate(epoch_results.values,
                        headers=columns,
                        tablefmt='fancy_grid'))

```

Val Loss	Train Loss	Val Accuracy	Train Accuracy	LSTM cells	Dropout
0.0829701	0.0447081	0.971044	0.982567	64	0.3
0.0801187	0.0529833	0.96816	0.979529	64	0.5
0.0936313	0.0711846	0.967255	0.976356	64	0.7
0.107525	0.0581586	0.96488	0.978963	32	0.5
0.113176	0.0646532	0.961769	0.97792	32	0.3
0.132993	0.080815	0.960355	0.97257	32	0.7
0.139463	0.134935	0.956566	0.948699	16	0.3
0.136943	0.116637	0.954021	0.95287	16	0.5
0.175511	0.182294	0.919692	0.914558	16	0.7

```

In [102]: r = talos.Reporting(t)

```

```

In [115]: r.best_params()

```

```

Out[115]: array([[3, '20', '64', 'binary_crossentropy', 'Adam', '16', '0.3', 0],
                 [7, '20', '64', 'binary_crossentropy', 'Adam', '16', '0.5', 1],
                 [0, '20', '64', 'binary_crossentropy', 'Adam', '16', '0.7', 2],
                 [2, '20', '32', 'binary_crossentropy', 'Adam', '16', '0.5', 3],
                 [8, '20', '32', 'binary_crossentropy', 'Adam', '16', '0.3', 4],
                 [6, '20', '32', 'binary_crossentropy', 'Adam', '16', '0.7', 5],
                 [5, '20', '16', 'binary_crossentropy', 'Adam', '16', '0.3', 6],
                 [1, '20', '16', 'binary_crossentropy', 'Adam', '16', '0.5', 7],
                 [4, '20', '16', 'binary_crossentropy', 'Adam', '16', '0.7', 8]],
                dtype=object)

```

```

In [142]: # from keras.callbacks import Callback
          # from IPython.display import clear_output

In [143]: # Code Courtesy : https://gist.github.com/stared/dfb4dfaf6d9a8501cd1cc8b8cb806d2e

          # class PlotLearning(Callback):
          #     def on_train_begin(self, logs={}):
          #         self.i = 0
          #         self.x = []
          #         self.losses = []
          #         self.val_losses = []
          #         self.acc = []
          #         self.val_acc = []
          #         self.fig = plt.figure()

          #         self.logs = []

          #     def on_epoch_end(self, epoch, logs={}):

          #         self.logs.append(logs)
          #         self.x.append(self.i)
          #         self.losses.append(logs.get('loss'))
          #         self.val_losses.append(logs.get('val_loss'))
          #         self.acc.append(logs.get('acc'))
          #         self.val_acc.append(logs.get('val_acc'))
          #         self.i += 1
          #         f, (ax1, ax2) = plt.subplots(1, 2, sharex=True)

          #         clear_output(wait=True)

          #         ax1.set_yscale('log')
          #         ax1.plot(self.x, self.losses, label="loss")
          #         ax1.plot(self.x, self.val_losses, label="val_loss")
          #         ax1.legend()

          #         ax2.plot(self.x, self.acc, label="accuracy")
          #         ax2.plot(self.x, self.val_acc, label="validation accuracy")
          #         ax2.legend()

          #         plt.show();

          # plot = PlotLearning()

```

#### 0.0.4 Randomised GridSearch with 2 layer GridSearch

```

In [149]: def get_best_lstm_model_2_layer(x_train, y_test, x_val, y_val, params):

          timesteps, data_point_size = X_train.shape[1:]

```

```

n_classes = y_train.shape[1]

model = Sequential()

model.add(LSTM(units=params['first_neuron'],
               input_shape=(timesteps, data_point_size),
               return_sequences=True))
model.add(Dropout(rate=params['dropout']))

model.add(LSTM(units=params['first_neuron']))
model.add(Dropout(rate=params['dropout']))

model.add(Dense(n_classes, activation='softmax'))

model.compile(optimizer=params['optimizer'],
              loss=params['losses'],
              metrics=['accuracy'])

history = model.fit(x_train, y_train,
                    validation_data=[x_val, y_val],
                    batch_size=params['batch_size'],
                    callbacks=[live()],
                    epochs=params['epochs'],
                    verbose=2)

return history, model

```

```


In [147]: params = {
    'optimizer': ['Adam'],
    'losses': ['binary_crossentropy'],
    'first_neuron': [64],
    'batch_size': [16],
    'epochs': [20],
    'dropout': [0.3, 0.5]
}

```

```

In [152]: t = talos.Scan(x=X_train, y=y_train,
                        x_val=X_test, y_val=y_test,
                        model=get_best_lstm_model_2_layer,
                        params=params, search_method='random',
                        shuffle=True, dataset_name='HAR_RAW',
                        experiment_no='2', seed=42)

```



100%|| 2/2 [2:21:02<00:00, 4233.94s/it]

Scan Finished!

In [153]: t.details

```
Out[153]: complete_time      02/02/19/20:30
          experiment_name    HAR_RAW_2
          grid_downsample    None
          random_method      uniform_mersenne
          reduce_loss        False
          reduction_interval  50
          reduction_method    None
          reduction_metric    val_acc
          reduction_threshold 0.2
          reduction_window    20
          x_shape             (7352, 128, 9)
          y_shape             (7352, 6)
          dtype: object
```

In [154]: print('We have fit {} number of models'.format(len(t.saved\_models)))

We have fit 2 number of models

```
In [155]: epoch_results = t.data.drop(['losses', 'optimizer',
                                         'epochs', 'batch_size', 'round_epochs'], axis=1)
          epoch_results['val_loss'] = epoch_results.val_loss.astype(float, inplace=True)
          epoch_results['loss'] = epoch_results.loss.astype(float, inplace=True)
          epoch_results['val_acc'] = epoch_results.val_acc.astype(float, inplace=True)
          epoch_results['acc'] = epoch_results.acc.astype(float, inplace=True)
          epoch_results.sort_values(by='val_acc', ascending=False, inplace=True)
```



```

columns = ['Val Loss', 'Train Loss',
           'Val Accuracy', 'Train Accuracy',
           'LSTM cells', 'Dropout']

In [156]: print(tabulate(epoch_results.values,
                        headers=columns,
                        tablefmt='fancy_grid'))

    Val Loss    Train Loss    Val Accuracy    Train Accuracy    LSTM cells    Dropout
-----
0.071176      0.037237      0.971553      0.983678          64           0.3
0.0857582     0.0404991     0.969743     0.984313          64           0.5

In [157]: r = talos.Reporting(t)

In [158]: r.best_params()

Out[158]: array([[ '20', '64', 'binary_crossentropy', 'Adam', '16', '0.3', 0],
                 [ '20', '64', 'binary_crossentropy', 'Adam', '16', '0.5', 1]],
                dtype=object)

```

### 0.0.5 Conclusion

- Talos is a wonderful library, a little buggy and difficult to get it perfect but reports many things and feels so native to GridSearchCV
- Dropouts are fitting the problem without much overfitting but this problem was computationally very expensive even on GPU
- Even 20 epochs seems more for the LSTMs. The problem is converging very soon within 15 epochs reminds me the power of LSTMs
- With LSTMs we are able to take the maximum validation accuracy to staggering 97.2% which amazing thinking the fact that our Domain expert features are absent here. Means a person without having any domain expertise also can do wonder here.
- Building 2 layer LSTM structure didn't do much good but it took more time. Hence it seems better to stick to 1 layer structure and looking at the number of training points it seems reasonable.
- One problems with talos is that it can parallelize a single model fit on GPU but can not parallelize all the model fits from GridSearch.